

Exploring the weather dataset

ANALYZING POLICE ACTIVITY WITH PANDAS



Kevin Markham


Founder, Data School

Introduction to the dataset



STATION DETAILS	
Name	PROVIDENCE, RI US
Network:ID	GHCND:USW00014765
Latitude/Longitude	41.7225°, -71.4325°
Elevation	16.8 m

PERIOD OF RECORD	
Start Date ¹	1942-08-01
End Date ¹	2018-04-21
Data Coverage ²	93%



```
weather = pd.read_csv('weather.csv')
weather.head(3)
```

	STATION	DATE	TAVG	TMIN	TMAX	AWND	WSF2	WT01	WT02	
0	USW00014765	2005-01-01	44.0	35	53	8.95	25.1	1.0	NaN	
1	USW00014765	2005-01-02	36.0	28	44	9.40	14.1	NaN	NaN	
2	USW00014765	2005-01-03	49.0	44	53	6.93	17.0	1.0	NaN	
	...									
	WT11	WT13	WT14	WT15	WT16	WT17	WT18	WT19	WT21	WT22
0	...	NaN	1.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1	...	NaN	NaN	NaN	NaN	1.0	NaN	1.0	NaN	NaN
2	...	NaN	1.0	NaN	NaN	1.0	NaN	NaN	NaN	NaN

- TAVG , TMIN , TMAX : Temperature
- AWND , WSF2 : Wind speed
- WT01 ... WT22 : Bad weather conditions

Examining the wind speed

```
weather[['AWND', 'WSF2']].head()
```

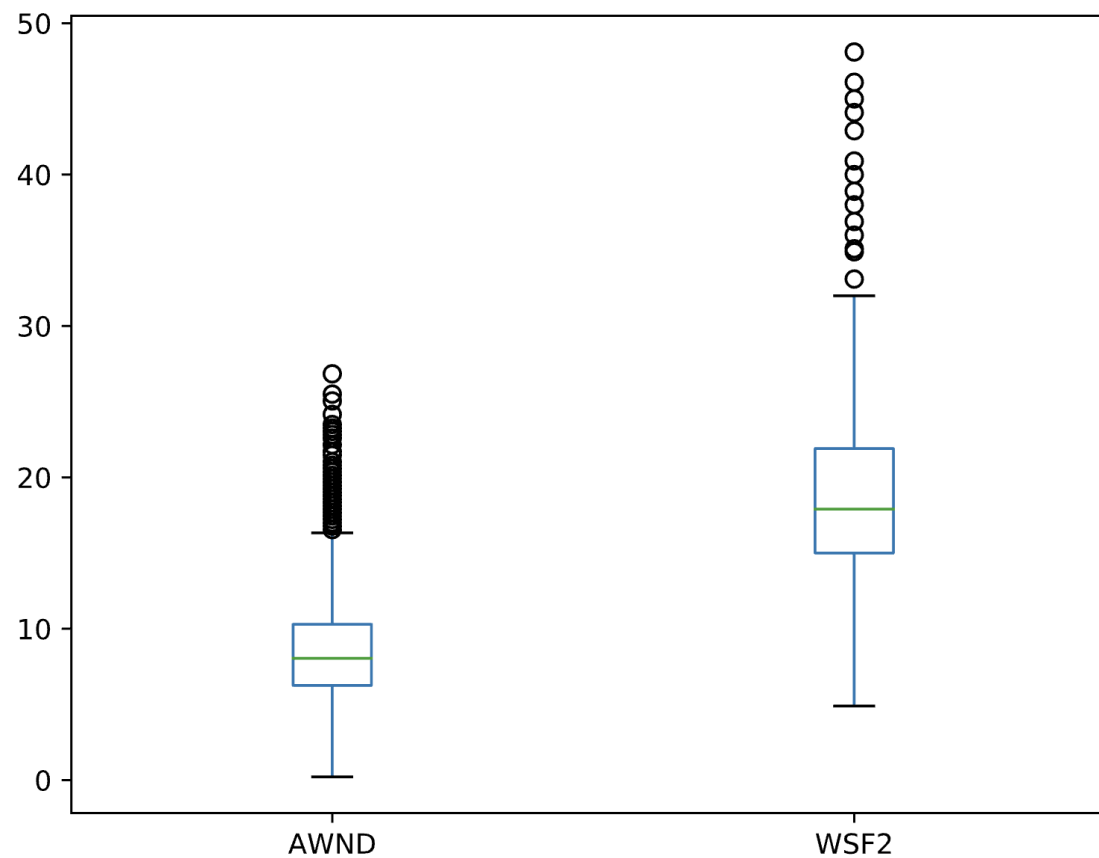
	AWND	WSF2
0	8.95	25.1
1	9.40	14.1
2	6.93	17.0
3	6.93	16.1
4	7.83	17.0

```
weather[['AWND', 'WSF2']].describe()
```

	AWND	WSF2
count	4017.000000	4017.000000
mean	8.593707	19.274782
std	3.364601	5.623866
min	0.220000	4.900000
25%	6.260000	15.000000
50%	8.050000	17.900000
75%	10.290000	21.900000
max	26.840000	48.100000

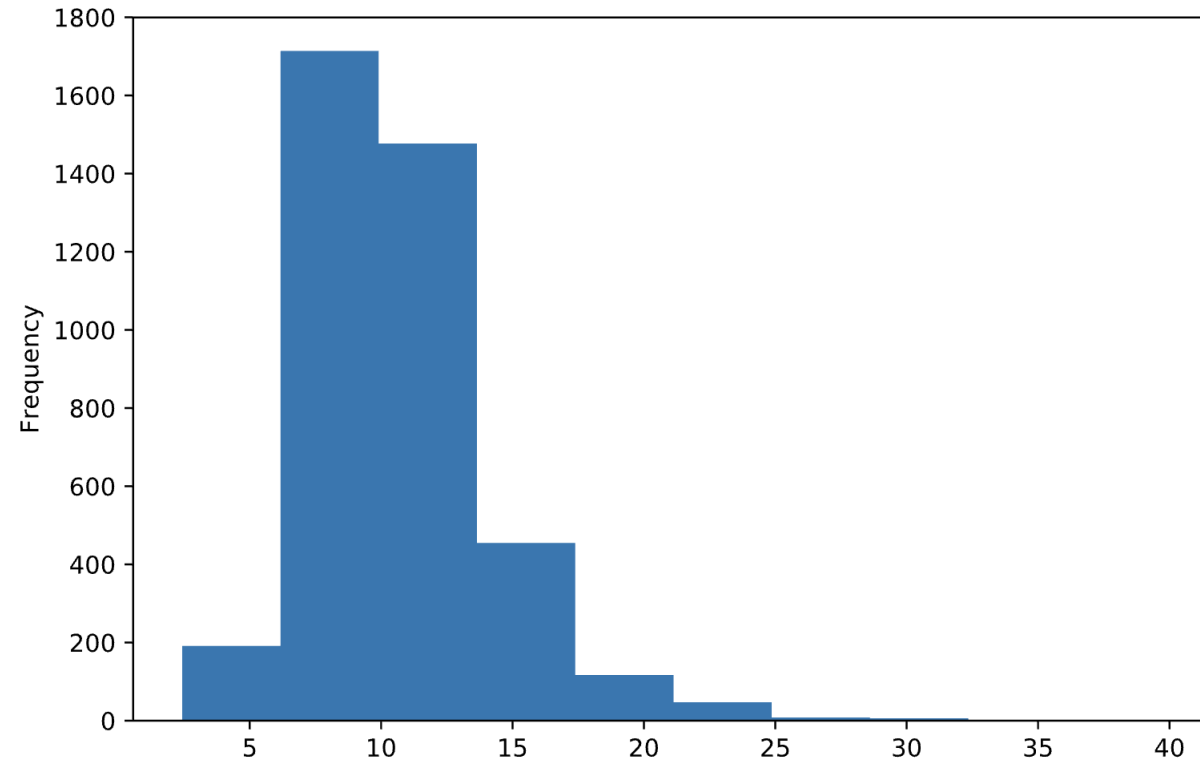
Creating a box plot

```
weather[['AWND', 'WSF2']].plot(kind='box')  
plt.show()
```



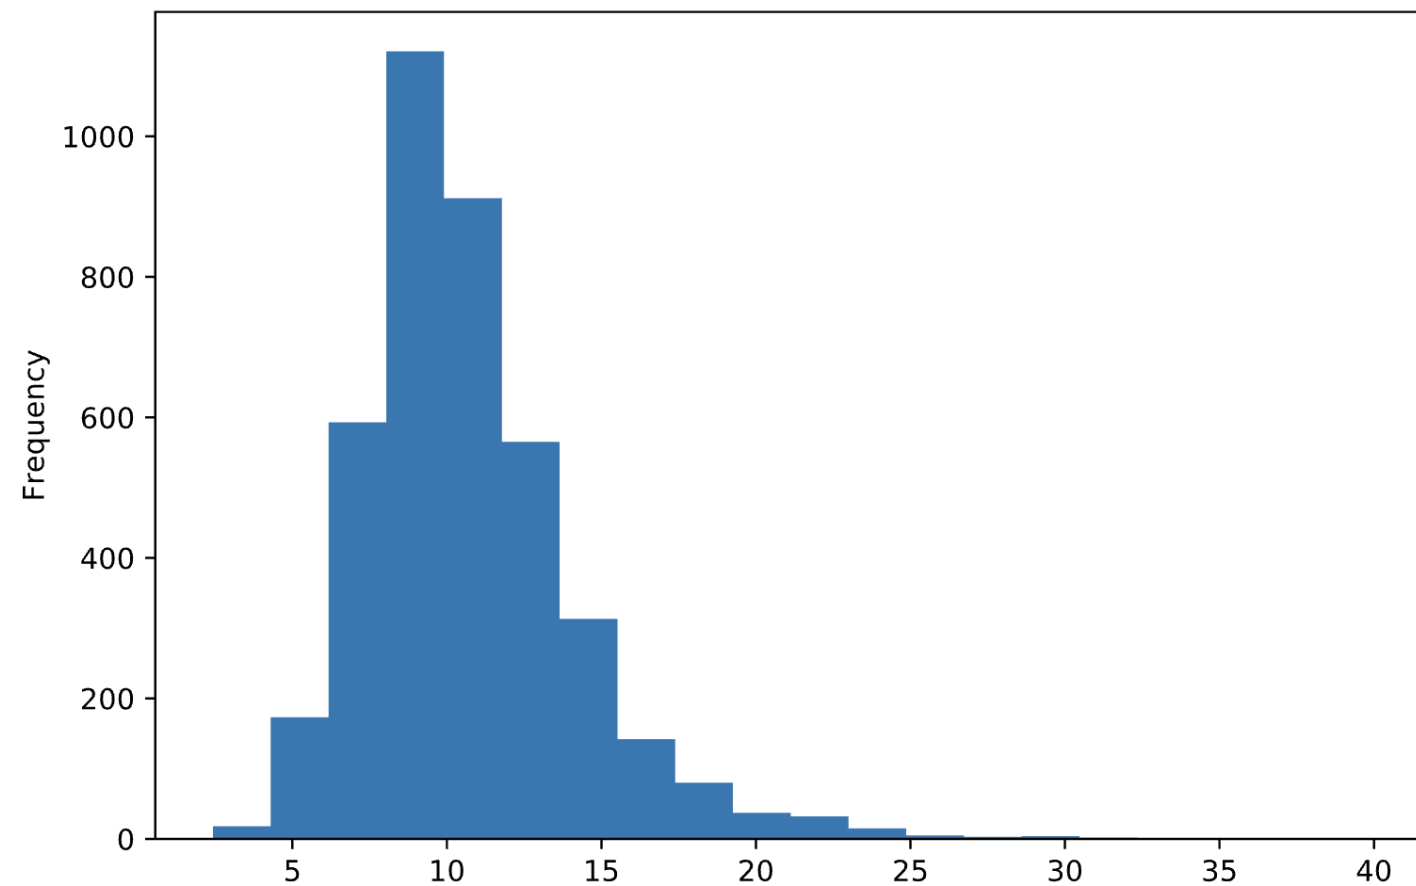
Creating a histogram (1)

```
weather['WDIFF'] = weather.WSF2 - weather.AWND  
weather.WDIFF.plot(kind='hist')  
plt.show()
```



Creating a histogram (2)

```
weather.WDIFF.plot(kind='hist', bins=20)  
plt.show()
```



Let's practice!

ANALYZING POLICE ACTIVITY WITH PANDAS

Categorizing the weather

ANALYZING POLICE ACTIVITY WITH PANDAS



Kevin Markham
Founder, Data School

Selecting a DataFrame slice (1)

```
weather.shape
```

```
(4017, 28)
```

```
weather.columns
```

```
Index(['STATION', 'DATE', 'TAVG', 'TMIN', 'TMAX', 'AWND',  
      'WSF2', 'WT01', 'WT02', 'WT03', 'WT04', 'WT05',  
      'WT06', 'WT07', 'WT08', 'WT09', 'WT10', 'WT11',  
      'WT13', 'WT14', 'WT15', 'WT16', 'WT17', 'WT18',  
      'WT19', 'WT21', 'WT22', 'TDIFF'], dtype='object')
```

Selecting a DataFrame slice (2)

```
temp = weather.loc[:, 'TAVG':'TMAX']
```

```
temp.shape
```

```
(4017, 3)
```

```
temp.columns
```

```
Index(['TAVG', 'TMIN', 'TMAX'], dtype='object')
```

DataFrame operations

```
temp.head()
```

	TAVG	TMIN	TMAX
0	44.0	35	53
1	36.0	28	44
2	49.0	44	53
3	42.0	39	45
4	36.0	28	43

```
temp.sum(axis='columns').head()
```

0	132.0
1	108.0
2	146.0
3	126.0
4	107.0

```
temp.sum()
```

TAVG	63884.0
TMIN	174677.0
TMAX	246116.0

Mapping one set of values to another

```
ri.stop_duration.unique()
```

```
array(['0-15 Min', '16-30 Min', '30+ Min'], dtype=object)
```

```
mapping = {'0-15 Min': 'short',  
          '16-30 Min': 'medium',  
          '30+ Min': 'long'}  
ri['stop_length'] = ri.stop_duration.map(mapping)  
ri.stop_length.dtype
```

```
dtype('O')
```

Changing data type from object to category (1)

```
ri.stop_length.unique()
```

```
array(['short', 'medium', 'long'], dtype=object)
```

- Category type stores the data more efficiently
- Allows you to specify a logical order for the categories

```
ri.stop_length.memory_usage(deep=True)
```

```
8689481
```

Changing data type from object to category (2)

```
cats = ['short', 'medium', 'long']
```

```
ri['stop_length'] = ri.stop_length.astype('category',  
                                          ordered=True,  
                                          categories=cats)
```

```
ri.stop_length.memory_usage(deep=True)
```

```
3400602
```

Using ordered categories (1)

```
ri.stop_length.head()
```

```
stop_datetime
2005-01-04 12:55:00    short
2005-01-23 23:15:00    short
2005-02-17 04:15:00    short
2005-02-20 17:15:00    medium
2005-02-24 01:20:00    short
Name: stop_length, dtype: category
Categories (3, object): [short < medium < long]
```


Using ordered categories (2)

```
ri[ri.stop_length > 'short'].shape
```

```
(16959, 16)
```

```
ri.groupby('stop_length').is_arrested.mean()
```

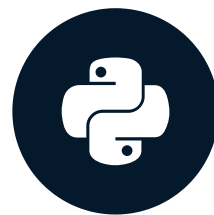
```
stop_length
short      0.013654
medium     0.093595
long       0.261572
Name: is_arrested, dtype: float64
```

Let's practice!

ANALYZING POLICE ACTIVITY WITH PANDAS

Merging datasets

ANALYZING POLICE ACTIVITY WITH PANDAS



Kevin Markham

Founder, Data School

```
apple
```

```
          date  time  price
date_and_time
2018-02-14 09:30:00  2/14/18   9:30  163.04
2018-02-14 16:00:00  2/14/18  16:00  167.37
2018-02-15 09:30:00  2/15/18   9:30  169.79
2018-02-15 16:00:00  2/15/18  16:00  172.99
```

```
apple.reset_index(inplace=True)
```

```
apple
```

```
   date_and_time  date  time  price
0 2018-02-14 09:30:00  2/14/18   9:30  163.04
1 2018-02-14 16:00:00  2/14/18  16:00  167.37
2 2018-02-15 09:30:00  2/15/18   9:30  169.79
3 2018-02-15 16:00:00  2/15/18  16:00  172.99
```

Preparing the second DataFrame

```
high_low
```

	DATE	HIGH	LOW
0	2/14/18	167.54	162.88
1	2/15/18	173.09	169.00
2	2/16/18	174.82	171.77

```
high = high_low[['DATE', 'HIGH']]  
high
```

	DATE	HIGH
0	2/14/18	167.54
1	2/15/18	173.09
2	2/16/18	174.82

Merging the DataFrames

```
apple_high = pd.merge(left=apple, right=high,  
                       left_on='date', right_on='DATE',  
                       how='left')
```

- `left=apple` : Left DataFrame
- `right=high` : Right DataFrame
- `left_on='date'` : Key column in left DataFrame
- `right_on='DATE'` : Key column in right DataFrame
- `how='left'` : Type of join

apple_high

	date_and_time		date	time	price	DATE	HIGH
0	2018-02-14	09:30:00	2/14/18	9:30	163.04	2/14/18	167.54
1	2018-02-14	16:00:00	2/14/18	16:00	167.37	2/14/18	167.54
2	2018-02-15	09:30:00	2/15/18	9:30	169.79	2/15/18	173.09
3	2018-02-15	16:00:00	2/15/18	16:00	172.99	2/15/18	173.09

apple

	date_and_time		date	time	price
0	2018-02-14	09:30:00	2/14/18	9:30	163.04
1	2018-02-14	16:00:00	2/14/18	16:00	167.37
2	2018-02-15	09:30:00	2/15/18	9:30	169.79
3	2018-02-15	16:00:00	2/15/18	16:00	172.99

high

	DATE	HIGH
0	2/14/18	167.54
1	2/15/18	173.09
2	2/16/18	174.82

Setting the index

```
apple_high.set_index('date_and_time', inplace=True)
apple_high
```

	date	time	price	DATE	HIGH
date_and_time					
2018-02-14 09:30:00	2/14/18	9:30	163.04	2/14/18	167.54
2018-02-14 16:00:00	2/14/18	16:00	167.37	2/14/18	167.54
2018-02-15 09:30:00	2/15/18	9:30	169.79	2/15/18	173.09
2018-02-15 16:00:00	2/15/18	16:00	172.99	2/15/18	173.09

Let's practice!

ANALYZING POLICE ACTIVITY WITH PANDAS

Does weather affect the arrest rate?

ANALYZING POLICE ACTIVITY WITH PANDAS



Kevin Markham
Founder, Data School

Driver gender and vehicle searches

```
ri.search_conducted.mean()
```

```
0.0382153092354627
```

```
ri.groupby('driver_gender').search_conducted.mean()
```

```
driver_gender
F      0.019181
M      0.045426
```

```
ri.groupby(['violation', 'driver_gender']).search_conducted.mean()
```

violation	driver_gender	
Equipment	F	0.039984
	M	0.071496
Moving violation	F	0.039257
	M	0.061524
Other	F	0.041018
	M	0.046191
Registration/plates	F	0.054924
	M	0.108802
Seat belt	F	0.017301
	M	0.035119
Speeding	F	0.008309
	M	0.027885

```
search_rate = ri.groupby(['violation',  
                          'driver_gender']).search_conducted.mean()  
  
search_rate
```

```
violation      driver_gender  
Equipment      F            0.039984  
               M            0.071496  
Moving violation F            0.039257  
               M            0.061524  
...           ...           ...
```

```
type(search_rate)  
type(search_rate.index)
```

```
pandas.core.series.Series  
pandas.core.indexes.multi.MultiIndex
```

violation	driver_gender	
Equipment	F	0.039984
	M	0.071496
Moving violation	F	0.039257
	M	0.061524
...

```
search_rate.loc['Equipment']
```

```
driver_gender
F    0.039984
M    0.071496
```

```
search_rate.loc['Equipment', 'M']
```

```
0.07149643705463182
```

Converting a multi-indexed Series to a DataFrame

```
search_rate.unstack()
```

```
driver_gender          F          M
violation
Equipment      0.039984  0.071496
Moving violation  0.039257  0.061524
Other           0.041018  0.046191
...              ...          ...
```

```
type(search_rate.unstack())
```

```
pandas.core.frame.DataFrame
```

Converting a multi-indexed Series to a DataFrame

```
ri.pivot_table(index='violation',  
                columns='driver_gender',  
                values='search_conducted')
```

driver_gender	F	M
violation		
Equipment	0.039984	0.071496
Moving violation	0.039257	0.061524
Other	0.041018	0.046191
...

Let's practice!

ANALYZING POLICE ACTIVITY WITH PANDAS

Conclusion

ANALYZING POLICE ACTIVITY WITH PANDAS



Kevin Markham

Founder, Data School

Stanford Open Policing Project



- Download data: <https://openpolicing.stanford.edu/>

Thank you!

ANALYZING POLICE ACTIVITY WITH PANDAS