

Using aggregation functions over windows

TIME SERIES ANALYSIS IN SQL SERVER



Kevin Feasel
CTO, Envizage

Ranking functions

`ROW_NUMBER()`

Unique, ascending integer value starting from 1.

`RANK()`

Ascending integer value starting from 1. Can have ties. Can skip numbers.

`DENSE_RANK()`

Ascending integer value starting from 1. Can have ties. Will not skip numbers.

RunsScored
8
7
7
6
6
3

Calculating row numbers

```
SELECT
    s.RunsScored,
    ROW_NUMBER() OVER (
        ORDER BY s.RunsScored DESC
    ) AS rn
FROM dbo.Scores s
ORDER BY
    s.RunsScored DESC;
```

RunsScored	rn
8	1
7	2
7	3
6	4
6	5
3	6

Calculating ranks and dense ranks

```
SELECT
    s.RunsScored,
    RANK() OVER (
        ORDER BY s.RunsScored DESC
    ) AS rk,
    DENSE_RANK() OVER (
        ORDER BY s.RunsScored DESC
    ) AS dr
FROM dbo.Scores s
ORDER BY
    s.RunsScored DESC;
```

RunsScored	rk	dr
8	1	1
7	2	2
7	2	2
6	4	3
6	4	3
3	6	4

Partitions

```
SELECT
    s.Team,
    s.RunsScored,
    ROW_NUMBER() OVER (
        PARTITION BY s.Team
        ORDER BY s.RunsScored DESC
    ) AS rn
FROM dbo.Scores s
ORDER BY
    s.RunsScored DESC;
```

Team	RunsScored	rn
AZ	8	1
AZ	6	2
AZ	3	3
FLA	7	1
FLA	7	2
FLA	6	3

Aggregate functions

```
SELECT
    s.Team,
    s.RunsScored,
    MAX(s.RunsScored) OVER (
        PARTITION BY s.Team
    ) AS MaxRuns
FROM dbo.Scores s
ORDER BY
    s.RunsScored DESC;
```

Team	RunsScored	MaxRuns
AZ	8	8
AZ	6	8
AZ	3	8
FLA	7	7
FLA	7	7
FLA	6	7

Aggregations with empty windows

```
SELECT
    s.Team,
    s.RunsScored,
    MAX(s.RunsScored) OVER() AS MaxRuns
FROM dbo.Scores s
ORDER BY
    s.RunsScored DESC;
```

Team	RunsScored	MaxRuns
AZ	8	8
AZ	6	8
AZ	3	8
FLA	7	8
FLA	7	8
FLA	6	8

Let's practice!

TIME SERIES ANALYSIS IN SQL SERVER

Calculating running totals and moving averages

TIME SERIES ANALYSIS IN SQL SERVER

SQL

Kevin Feasel
CTO, Envizage

Calculating running totals

Team	Game	RunsScored
AZ	1	8
AZ	2	6
AZ	3	3
FLA	1	7
FLA	2	7
FLA	3	6

Team	Game	RunsScored	TotalRuns
AZ	1	8	8
AZ	2	6	14
AZ	3	3	17
FLA	1	7	7
FLA	2	7	14
FLA	3	6	20

Running totals

```
SELECT
  s.Team,
  s.Game,
  s.RunsScored,
  SUM(s.RunsScored) OVER (
    PARTITION BY s.Team
    ORDER BY s.Game ASC
    RANGE BETWEEN
      UNBOUNDED PRECEDING
      AND CURRENT ROW
  ) AS TotalRuns
FROM #Scores s;
```

- Team , Game , RunsScored columns
- SUM(s.RunsScored)
- OVER()
- PARTITION BY s.Team
- ORDER BY s.Game ASC
- RANGE BETWEEN
- UNBOUNDED PRECEDING
- AND CURRENT ROW

RANGE and ROWS

RANGE

- Specify a range of results
- "Duplicates" processed all at once
- Only supports `UNBOUNDED` and `CURRENT ROW`

ROWS

- Specify number of rows to include
- "Duplicates" processed a row at a time
- Supports `UNBOUNDED` , `CURRENT ROW` , and number of rows

Calculating moving averages

```
SELECT
  s.Team,
  s.Game,
  s.RunsScored,
  AVG(s.RunsScored) OVER (
    PARTITION BY s.Team
    ORDER BY s.Game ASC
    ROWS BETWEEN 1 PRECEDING
    AND CURRENT ROW
  ) AS AvgRuns
FROM #Scores s;
```

Team	Game	RunsScored	AvgRuns
AZ	1	8	8
AZ	2	6	7
AZ	3	3	4
FLA	1	7	7
FLA	2	7	7
FLA	3	6	6

Let's practice!

TIME SERIES ANALYSIS IN SQL SERVER

Working with LAG() and LEAD()

TIME SERIES ANALYSIS IN SQL SERVER

SQL

Kevin Feasel
CTO, Envizage

The LAG() window function

```
SELECT
    dsr.CustomerID,
    dsr.MonthStartDate,
    LAG(dsr.NumberOfVisits) OVER (PARTITION BY dsr.CustomerID ORDER BY dsr.MonthStartDate) AS Prior,
    dsr.NumberOfVisits
FROM dbo.DaySpaRollup dsr;
```

CustomerID	MonthStartDate	Prior	NumberOfVisits
1	2018-12-01	NULL	49
1	2019-01-01	49	117
1	2019-02-01	117	104

The LEAD() window function

```
SELECT
    dsr.CustomerID,
    dsr.MonthStartDate,
    dsr.NumberOfVisits,
    LEAD(dsr.NumberOfVisits) OVER (PARTITION BY dsr.CustomerID ORDER BY dsr.MonthStartDate) AS Next
FROM dbo.DaySpaRollup dsr;
```

CustomerID	MonthStartDate	NumberOfVisits	Next
1	2018-12-01	49	117
1	2019-01-01	117	104
1	2019-02-01	104	108

Specifying number of rows back

```
SELECT
    dsr.CustomerID,
    dsr.MonthStartDate,
    LAG(dsr.NumberOfVisits, 2) OVER (PARTITION BY dsr.CustomerID ORDER BY dsr.MonthStartDate) AS Prior2,
    LAG(dsr.NumberOfVisits, 1) OVER (PARTITION BY dsr.CustomerID ORDER BY dsr.MonthStartDate) AS Prior1,
    dsr.NumberOfVisits
FROM dbo.DaySpaRollup dsr;
```

CustomerID	MonthStartDate	Prior2	Prior	NumberOfVisits
1	2018-12-01	NULL	NULL	49
1	2019-01-01	NULL	49	117
1	2019-02-01	49	117	104

```
SELECT
    Date,
    LAG(Val, 1) OVER(ORDER BY DATE) AS PriorVal,
    Val
FROM t;
```

Date	PriorVal	Val
2019-01-01	NULL	3
2019-01-02	3	6
2019-01-03	6	4

```
SELECT
    Date,
    LAG(Val, 1) OVER(ORDER BY DATE) AS PriorVal,
    Val
FROM t
WHERE
    t.Date > '2019-01-02';
```

Date	PriorVal	Val
2019-01-03	NULL	4

Windows and filters and CTEs

```
WITH records AS (  
  SELECT  
    Date,  
    LAG(Val, 1) OVER(ORDER BY Date) AS PriorVal,  
    Val  
  FROM t  
)  
SELECT  
  r.Date,  
  r.PriorVal,  
  r.Val  
FROM records r  
WHERE  
  r.Date > '2019-01-02';
```

Date	PriorVal	Val
2019-01-03	6	4

Let's practice!

TIME SERIES ANALYSIS IN SQL SERVER

Finding maximum levels of overlap

TIME SERIES ANALYSIS IN SQL SERVER



Kevin Feasel
CTO, Envizage

Start with some data

StartTime	EndTime	ProductsOrdered
2019-07-08 14:35:00	2019-07-08 16:01:00	13
2019-07-08 15:35:00	2019-07-08 17:01:00	13
2019-07-08 16:35:00	2019-07-08 18:01:00	17
2019-07-08 17:35:00	2019-07-08 19:01:00	15
2019-07-08 17:55:00	2019-07-08 17:57:00	1
2019-07-08 20:35:00	2019-07-08 22:01:00	13

Reasoning through the problem

StartTime	EndTime	ProductsOrdered
2019-07-08 14:35:00	2019-07-08 16:01:00	13
2019-07-08 15:35:00	2019-07-08 17:01:00	13
2019-07-08 16:35:00	2019-07-08 18:01:00	17
2019-07-08 17:35:00	2019-07-08 19:01:00	15
2019-07-08 17:55:00	2019-07-08 17:57:00	1
2019-07-08 20:35:00	2019-07-08 22:01:00	13

Reasoning through the problem

StartTime	EndTime	ProductsOrdered
2019-07-08 14:35:00	2019-07-08 16:01:00	13
2019-07-08 15:35:00	2019-07-08 17:01:00	13
2019-07-08 16:35:00	2019-07-08 18:01:00	17
2019-07-08 17:35:00	2019-07-08 19:01:00	15
2019-07-08 17:55:00	2019-07-08 17:57:00	1
2019-07-08 20:35:00	2019-07-08 22:01:00	13

Algorithm, step 1

CTE StartStopPoints:

```
SELECT
    o.StartTime AS TimeUTC,
    1 AS EntryCount,
    ROW_NUMBER() OVER (ORDER BY o.StartTime) AS StartOrdinal
FROM #Orders o
UNION ALL
SELECT
    o.EndTime AS TimeUTC,
    -1 AS EntryCount,
    NULL AS StartOrdinal
FROM #Orders o
```

Algorithm, step 1

TimeUTC	EntryCount	StartOrdinal
14:35:00	1	1
15:35:00	1	2
16:35:00	1	3
17:35:00	1	4
17:55:00	1	5
20:35:00	1	6

TimeUTC	EntryCount	StartOrdinal
16:01:00	-1	NULL
17:01:00	-1	NULL
18:01:00	-1	NULL
19:01:00	-1	NULL
17:57:00	-1	NULL
22:01:00	-1	NULL

Algorithm, step 2

CTE StartStopOrder:

```
SELECT
    s.TimeUTC,
    s.EntryCount,
    s.StartOrdinal,
    ROW_NUMBER() OVER (ORDER BY TimeUTC, StartOrdinal) AS StartOrEndOrdinal
FROM StartStopPoints s
```

Algorithm, step 2

TimeUTC	EC	SO	StartEndOrdinal
14:35:00	1	1	1
15:35:00	1	2	2
16:01:00	-1	NULL	3
16:35:00	1	3	4
17:01:00	-1	NULL	5
17:35:00	1	4	6

TimeUTC	EC	SO	StartEndOrdinal
17:55:00	1	5	7
17:57:00	-1	NULL	8
18:01:00	-1	NULL	9
19:01:00	-1	NULL	10
20:35:00	1	6	11
22:01:00	-1	NULL	12

Algorithm, step 2

TimeUTC	EC	SO	StartEndOrdinal
14:35:00	1	1	1
15:35:00	1	2	2
16:01:00	-1	NULL	3
16:35:00	1	3	4
17:01:00	-1	NULL	5
17:35:00	1	4	6

TimeUTC	EC	SO	StartEndOrdinal
17:55:00	1	5	7
17:57:00	-1	NULL	8
18:01:00	-1	NULL	9
19:01:00	-1	NULL	10
20:35:00	1	6	11
22:01:00	-1	NULL	12

Algorithm, step 3

TimeUTC	StartOrdinal	StartEndOrdinal	Calc	Result
14:35:00	1	1	$(2*1) - 1$	1
15:35:00	2	2	$(2*2) - 2$	2
16:01:00	NULL	3	NULL	NULL
16:35:00	3	4	$(2*3) - 4$	2
17:01:00	NULL	5	NULL	NULL
17:35:00	4	6	$(2*4) - 6$	2

Algorithm, step 3

```
SELECT  
    MAX(2 * s.StartOrdinal - s.StartOrEndOrdinal) AS MaxConcurrentVisitors  
FROM StartStopOrder s  
WHERE s.EntryCount = 1;
```

MaxConcurrentVisitors

3

Let's practice!

TIME SERIES ANALYSIS IN SQL SERVER

Wrapping up

TIME SERIES ANALYSIS IN SQL SERVER



Kevin Feasel
CTO, Envizage

Working with dates

- Combine `DATEADD()` and `DATEDIFF()` to round dates and times.
- Format with `CAST()` and `CONVERT()` when performance matters. `FORMAT()` is useful but slow.
- Calendar tables are a valuable asset.

Building dates

- `CAST()` , `CONVERT()` , and `PARSE()` can all turn strings into dates.
- Use `TRY_CAST()` , `TRY_CONVERT()` , and `TRY_PARSE()` for safe date conversions.
- `SWITCHOFFSET()` and `TODATETIMEOFFSET()` are useful functions for working with offsets.

Time-based aggregates

- Aggregate functions include `COUNT()` , `MIN()` , `MAX()` , and `SUM()` .
- Statistical aggregate functions include `AVG()` , `STDEV()` , `VAR()` , `STDEVP()` , and `VARP()` .
- `ROLLUP` , `CUBE` , and `GROUPING SETS` allow you to refine your aggregations.

Common (and uncommon) time series problems

- Windows work over ranking functions (`ROW_NUMBER()` , `RANK()` , `DENSE_RANK()` , and `NTILE()`) as well as aggregate functions (including statistical functions).
- Running totals are a use of `SUM()` over a window.
- Moving averages are a use of `AVG()` over a window.
- `LAG()` and `LEAD()` let us peek backward and forward in time.
- Pivot and transform date data to calculate concurrency.

Grazie!

TIME SERIES ANALYSIS IN SQL SERVER