Probability mass functions

EXPLORATORY DATA ANALYSIS IN PYTHON

Allen Downey
Professor, Olin College

datacamp

GSS

- Annual sample of U.S. population.
- Asks about demographics, social and political beliefs.
- Widely used by policy makers and researchers.

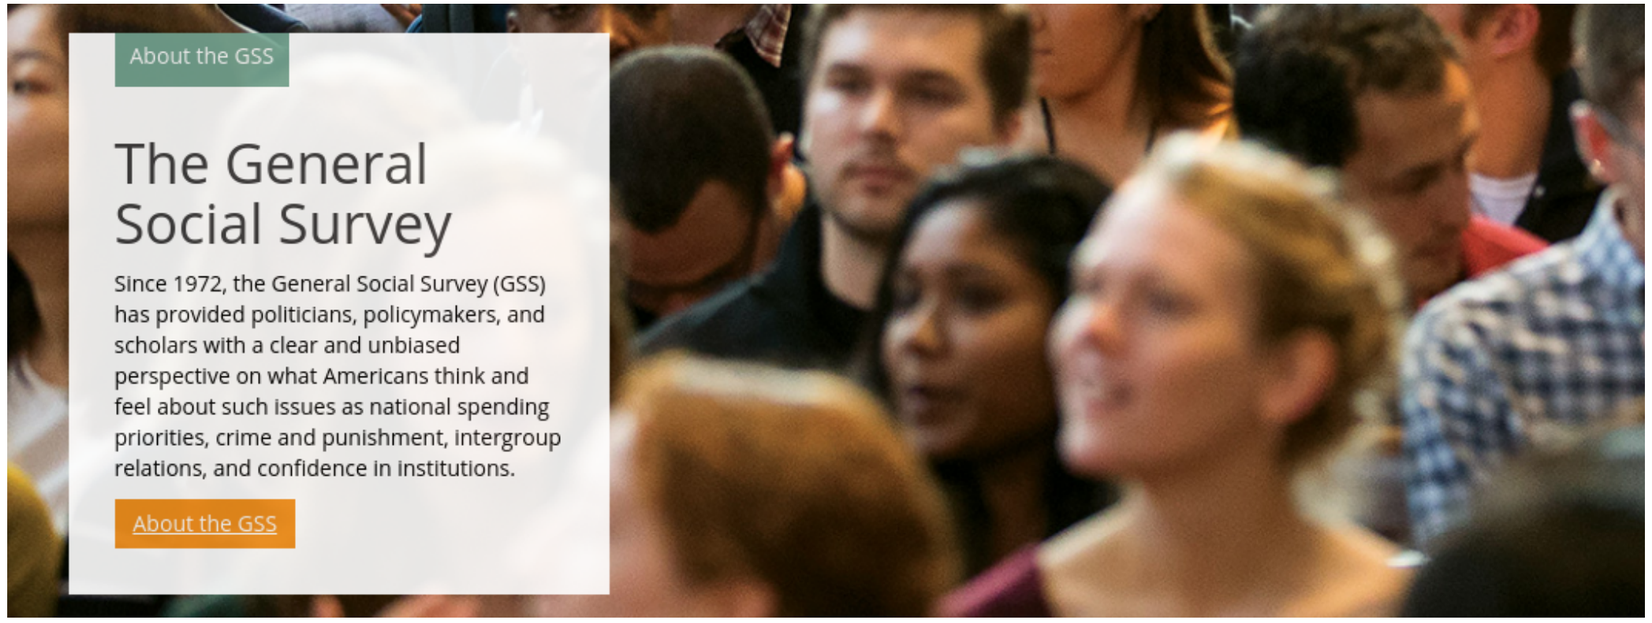

Read the data

```
gss = pd.read_hdf('gss.hdf5', 'gss')

gss.head()
```

	year	sex	age	cohort	race	educ	realinc	wtssall
0	1972	1	26.0	1946.0	1	18.0	13537.0	0.8893
1	1972	2	38.0	1934.0	1	12.0	18951.0	0.4446
2	1972	1	57.0	1915.0	1	12.0	30458.0	1.3339
3	1972	2	61.0	1911.0	1	14.0	37226.0	0.8893
4	1972	1	59.0	1913.0	1	12.0	30458.0	0.8893

```
educ = gss['educ']
plt.hist(educ.dropna(), label='educ')
plt.show()
```

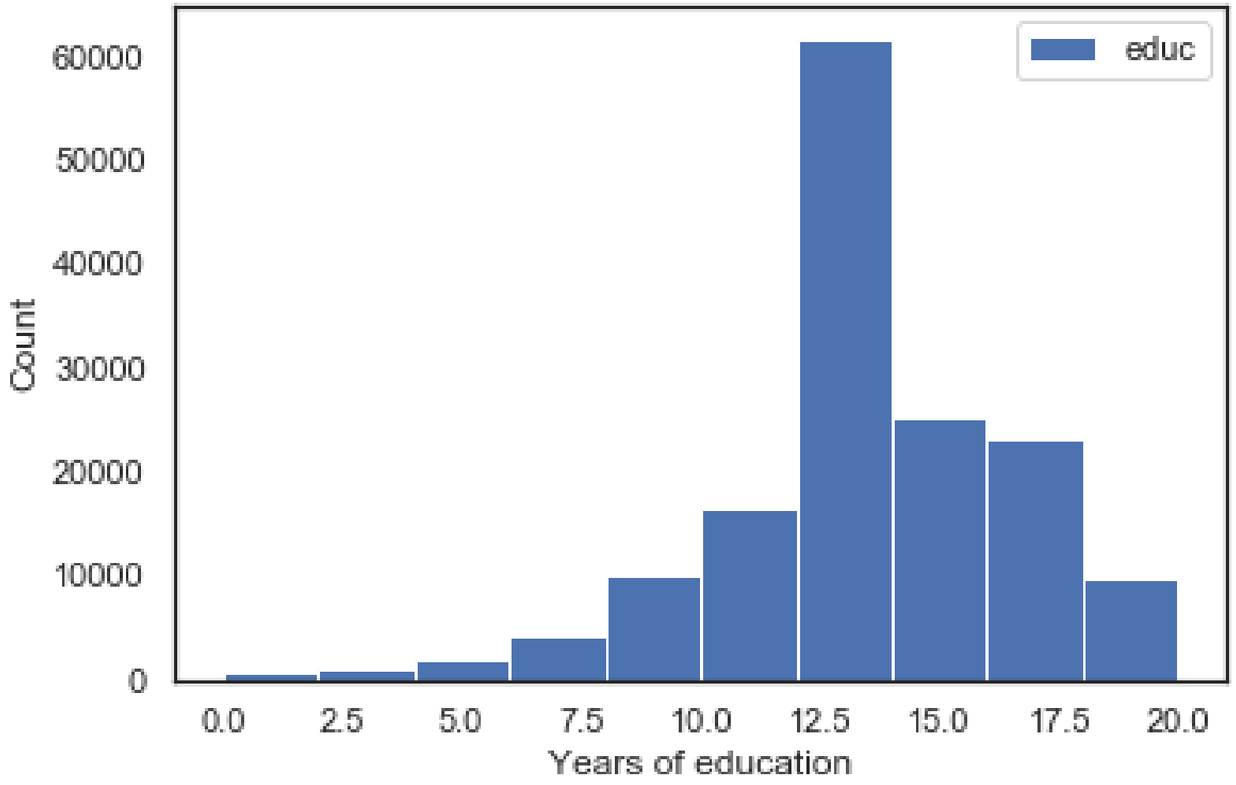

PMF

```
pmf_educ = Pmf(educ, normalize=False)
pmf_educ.head()
```

```
0.0    566
1.0    118
2.0    292
3.0    686
4.0    746
Name: educ, dtype: int64
```

PMF

```
pmf_educ[12]
```

```
47689
```

```
pmf_educ = Pmf(educ, normalize=True)


pmf_educ.head()
```

```
0.0    0.003663
1.0    0.000764
2.0    0.001890
3.0    0.004440
4.0    0.004828
Name: educ, dtype: int64
```

```
pmf_educ[12]
```

```
0.30863869940587907
```

```
pmf_educ.bar(label='educ')
plt.xlabel('Years of education')
plt.ylabel('PMF')
plt.show()
```

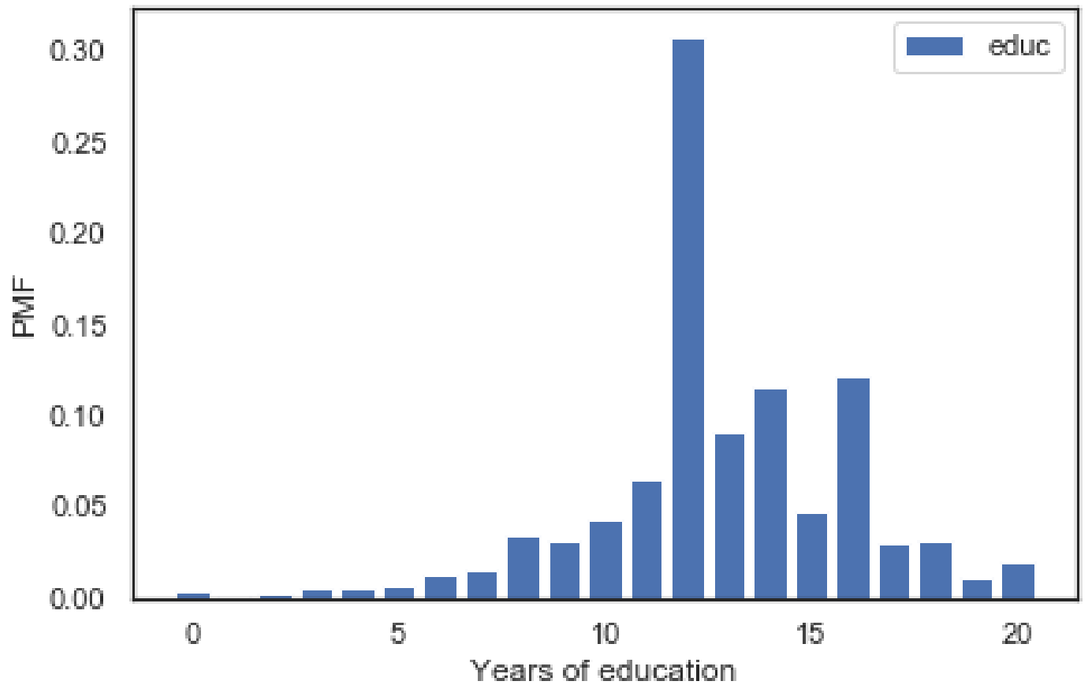

Histogram vs. PMF

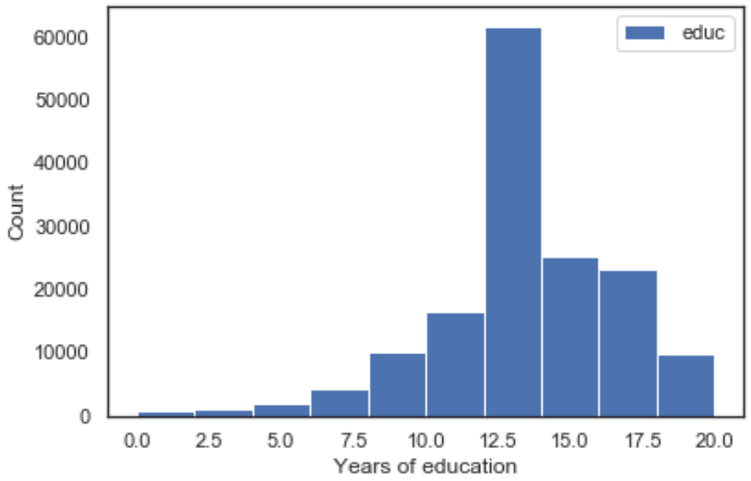

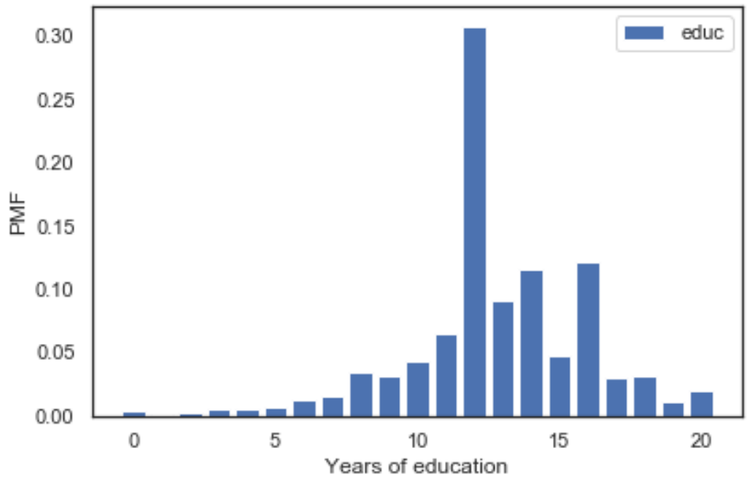

Let's make some PMFs!

EXPLORATORY DATA ANALYSIS IN PYTHON

Cumulative distribution functions

EXPLORATORY DATA ANALYSIS IN PYTHON

Allen Downey
Professor, Olin College

datacamp

From PMF to CDF

If you draw a random element from a distribution:

- PMF (Probability Mass Function) is the probability that you get exactly x
- CDF (Cumulative Distribution Function) is the probability that you get a value <= x

for a given value of x.

Example

PMF of {1, 2, 2, 3, 5}

PMF(1) = 1/5

PMF(2) = 2/5

PMF(3) = 1/5

PMF(5) = 1/5

CDF is the cumulative sum of the PMF.

CDF(1) = 1/5

CDF(2) = 3/5

CDF(3) = 4/5

CDF(5) = 1

```
cdf = Cdf(gss['age'])
cdf.plot()
plt.xlabel('Age')
plt.ylabel('CDF')
plt.show()
```

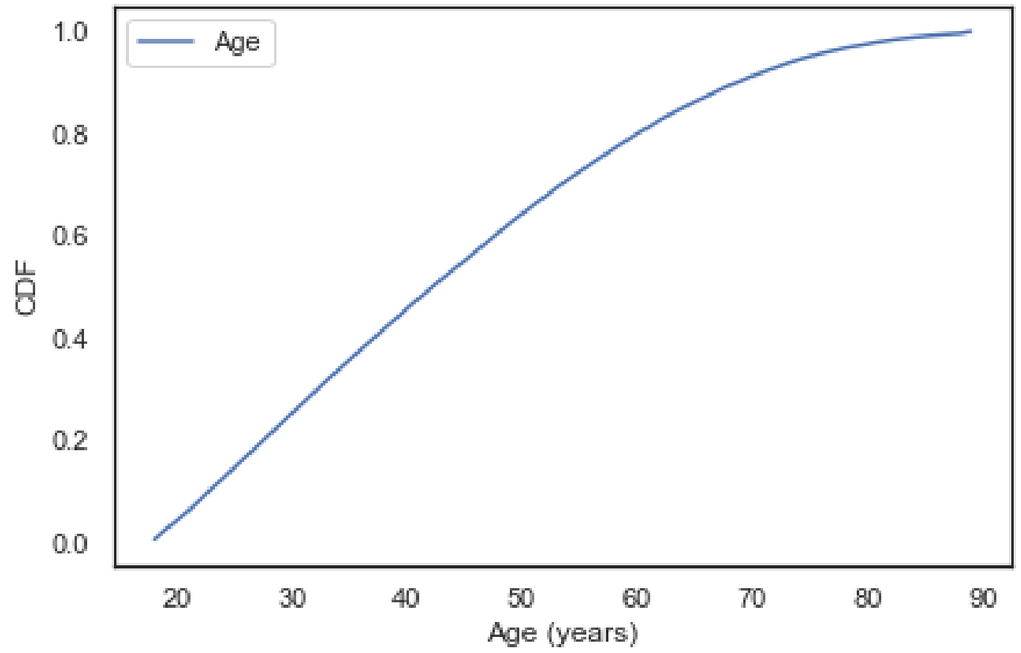

Evaluating the CDF

```
q = 51
p = cdf(q)
print(p)
```

```
0.66
```

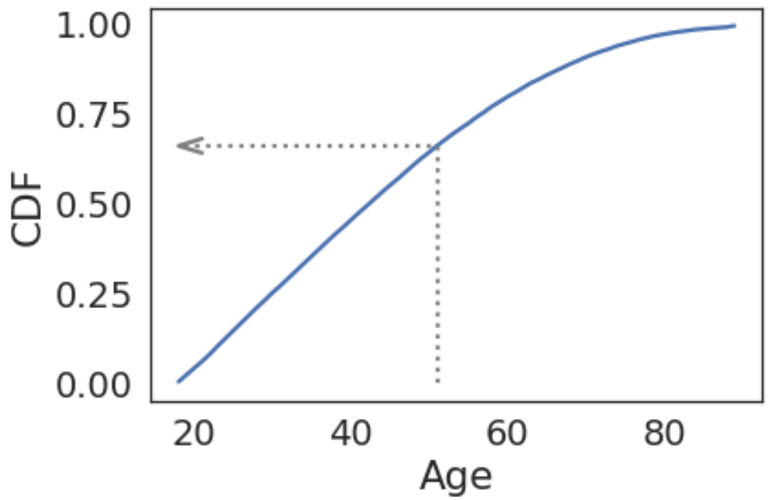

Evaluating the inverse CDF

```
p = 0.25
q = cdf.inverse(p)
print(q)
```

```
30
```

```
p = 0.75
q = cdf.inverse(p)
print(q)
```

```
57
```

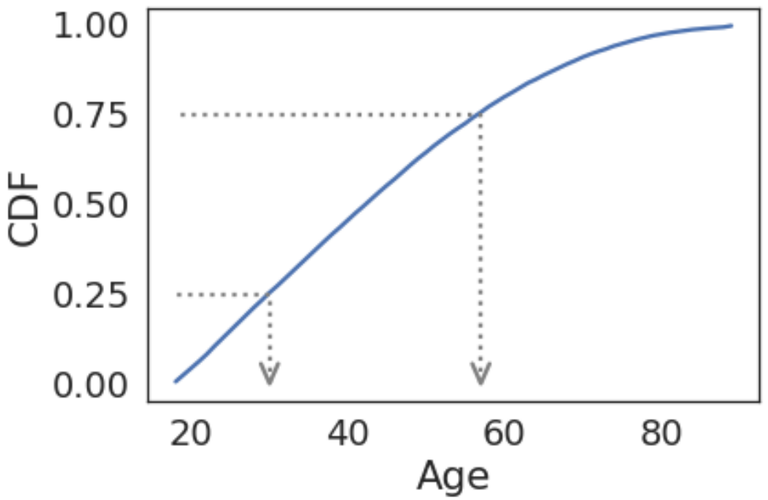

Let's practice!

EXPLORATORY DATA ANALYSIS IN PYTHON

Comparing distributions

EXPLORATORY DATA ANALYSIS IN PYTHON

Allen Downey
Professor, Olin College

datacamp

Multiple PMFs

```
male = gss['sex'] == 1
age = gss['age']
male_age = age[male]
female_age = age[~male]
Pmf(male_age).plot(label='Male')
Pmf(female_age).plot(label='Female')
plt.xlabel('Age (years)')
plt.ylabel('Count')
plt.show()
```

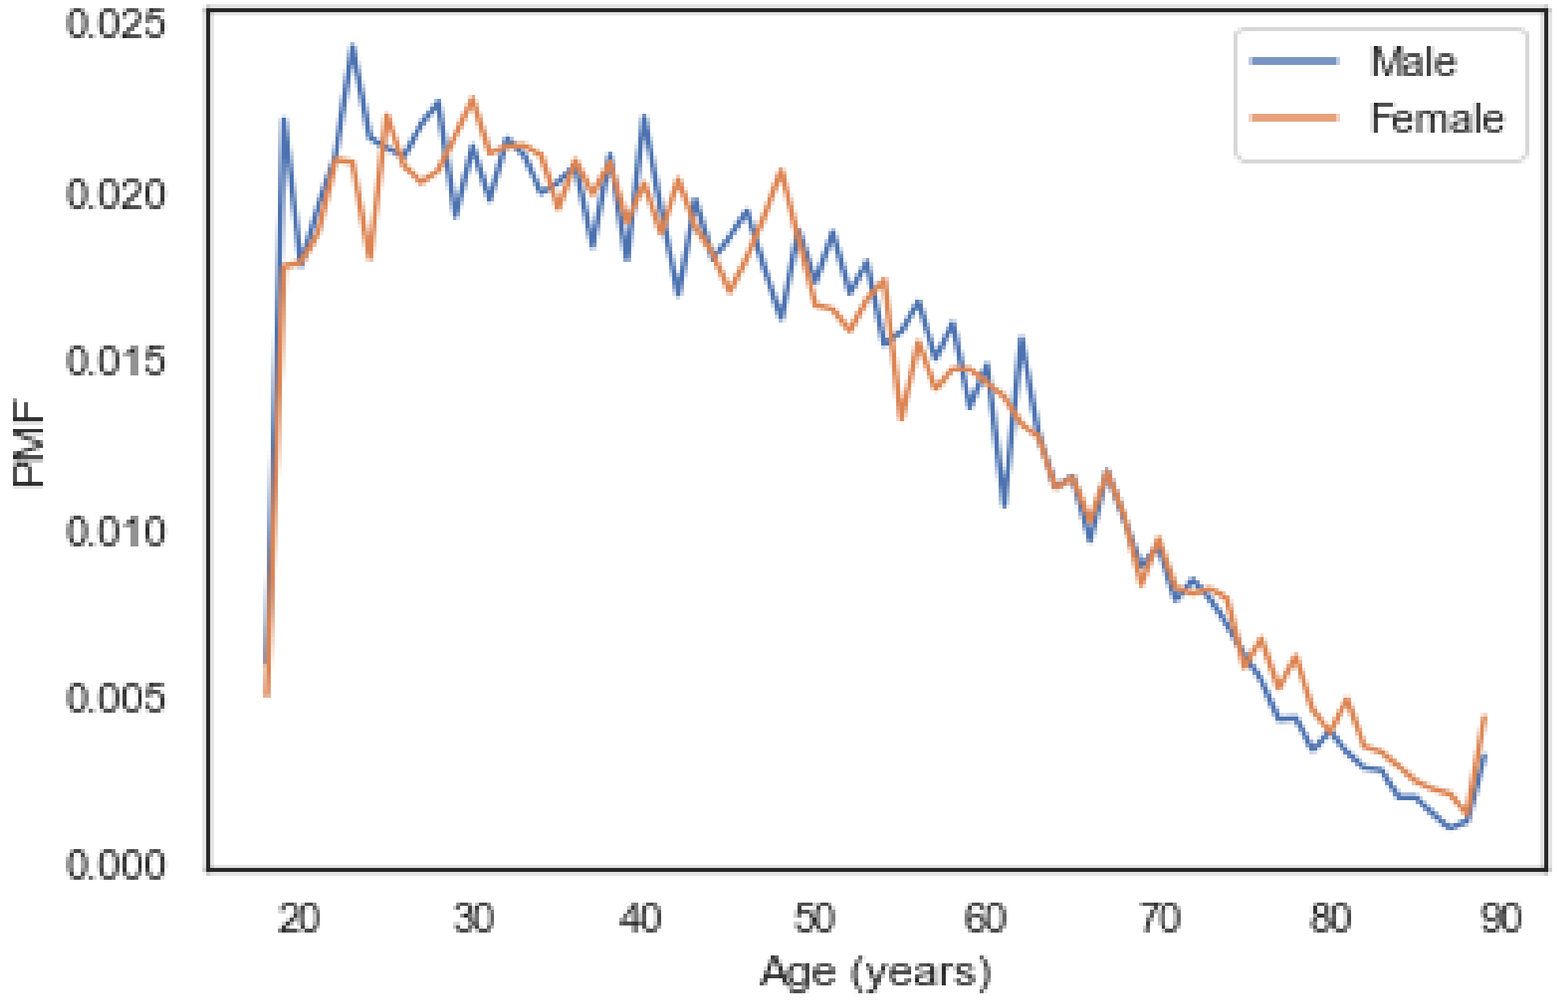
0.025
0.020
0.015
0.010
0.005
0.000
PMF
20
30
40
50
60
70
80
90
Age (years)
Male
Female

Multiple CDFs

```
Cdf(male_age).plot(label='Male')
Cdf(female_age).plot(label='Female')


plt.xlabel('Age (years)')
plt.ylabel('Count')
plt.show()
```

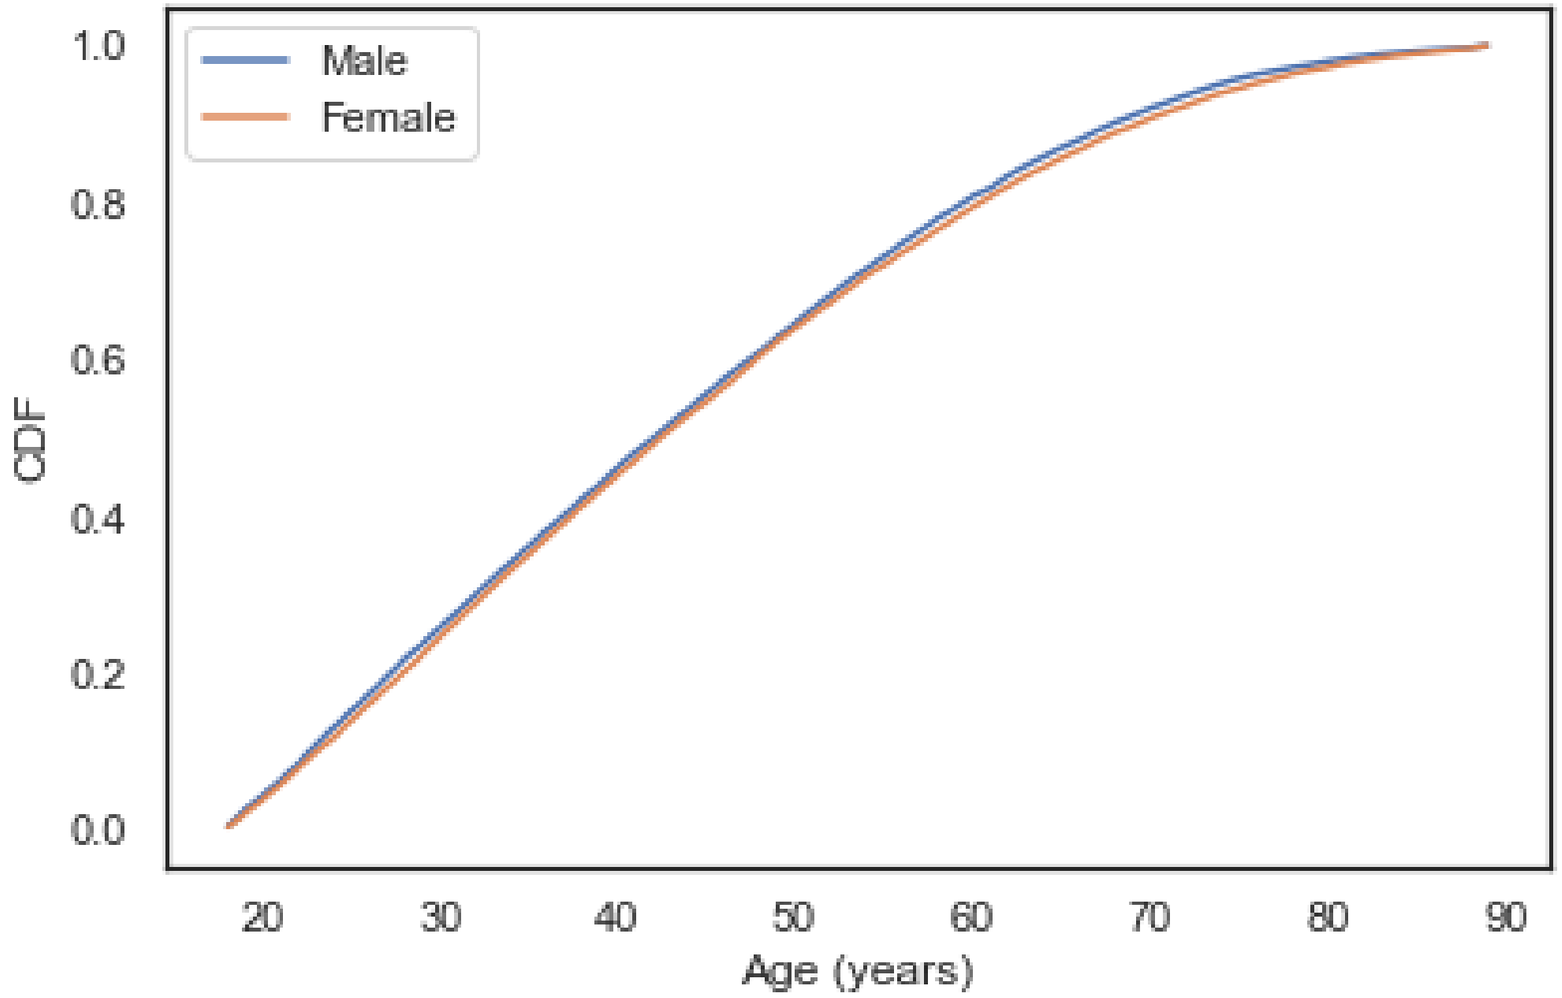
Male
Female
CDF
1.0
0.8
0.6
0.4
0.2
0.0
20
30
40
50
60
70
80
90
Age (years)

Income distribution

```
income = gss['realinc']
pre95 = gss['year'] < 1995
Pmf(income[pre95]).plot(label='Before 1995')
Pmf(income[~pre95]).plot(label='After 1995')
plt.xlabel('Income (1986 USD)')
plt.ylabel('PMF')
plt.show()
```

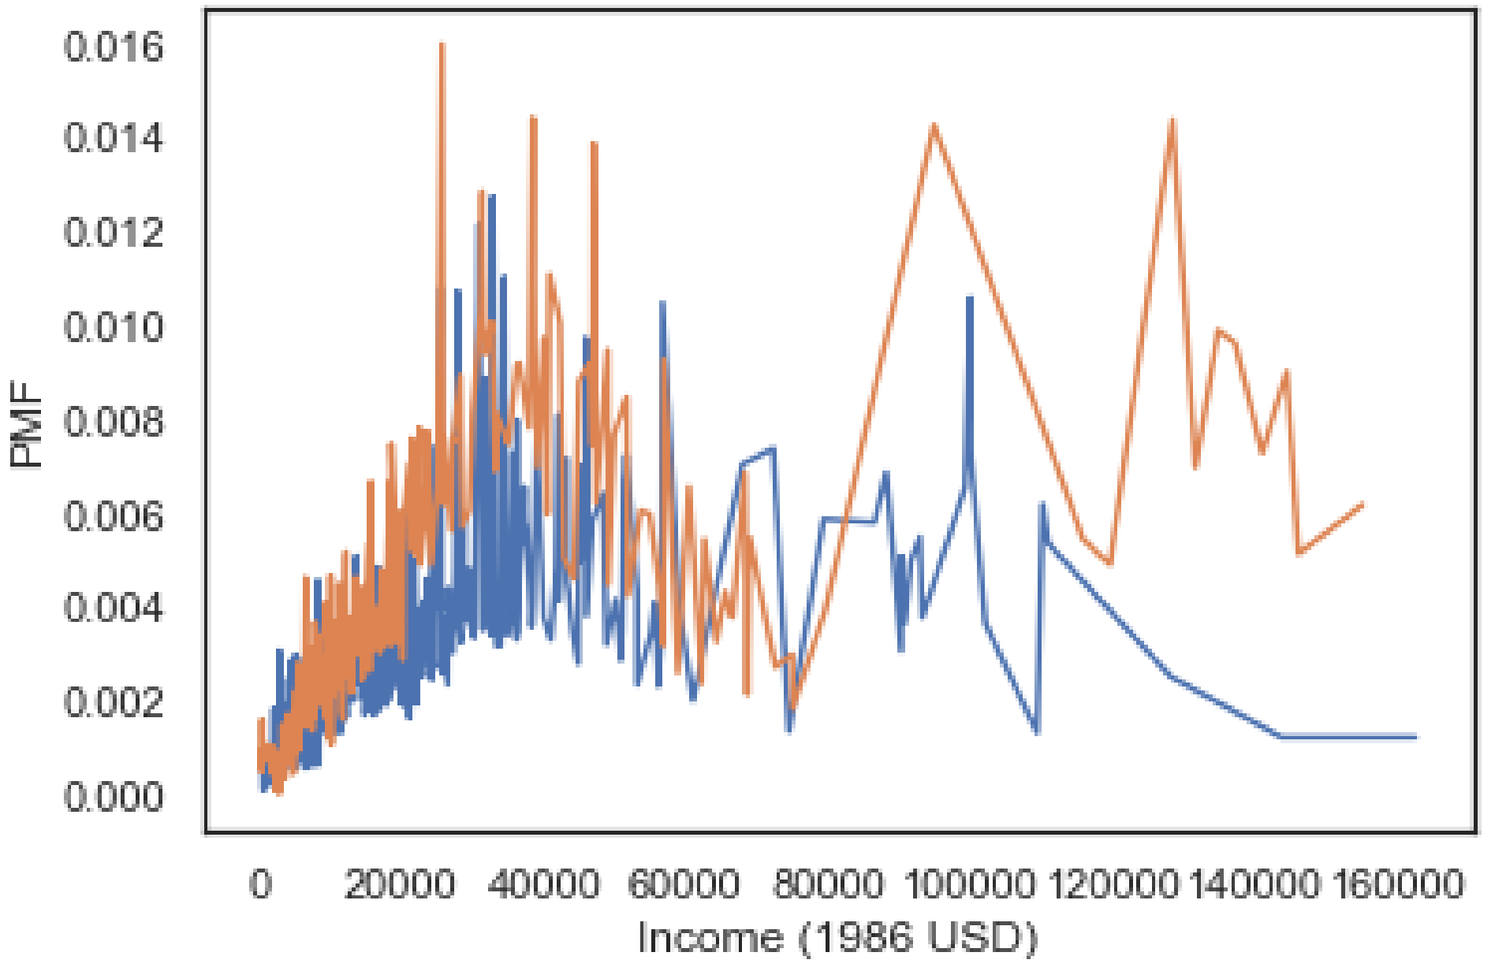
0.016
0.014
0.012
0.010
0.008
0.006
0.004
0.002
0.000
PMF
0
20000
40000
60000
80000
100000
120000
140000
160000
Income (1986 USD)

Income CDFs

```
Cdf(income[pre95]).plot(label='Before 1995')
Cdf(income[~pre95]).plot(label='After 1995')
```

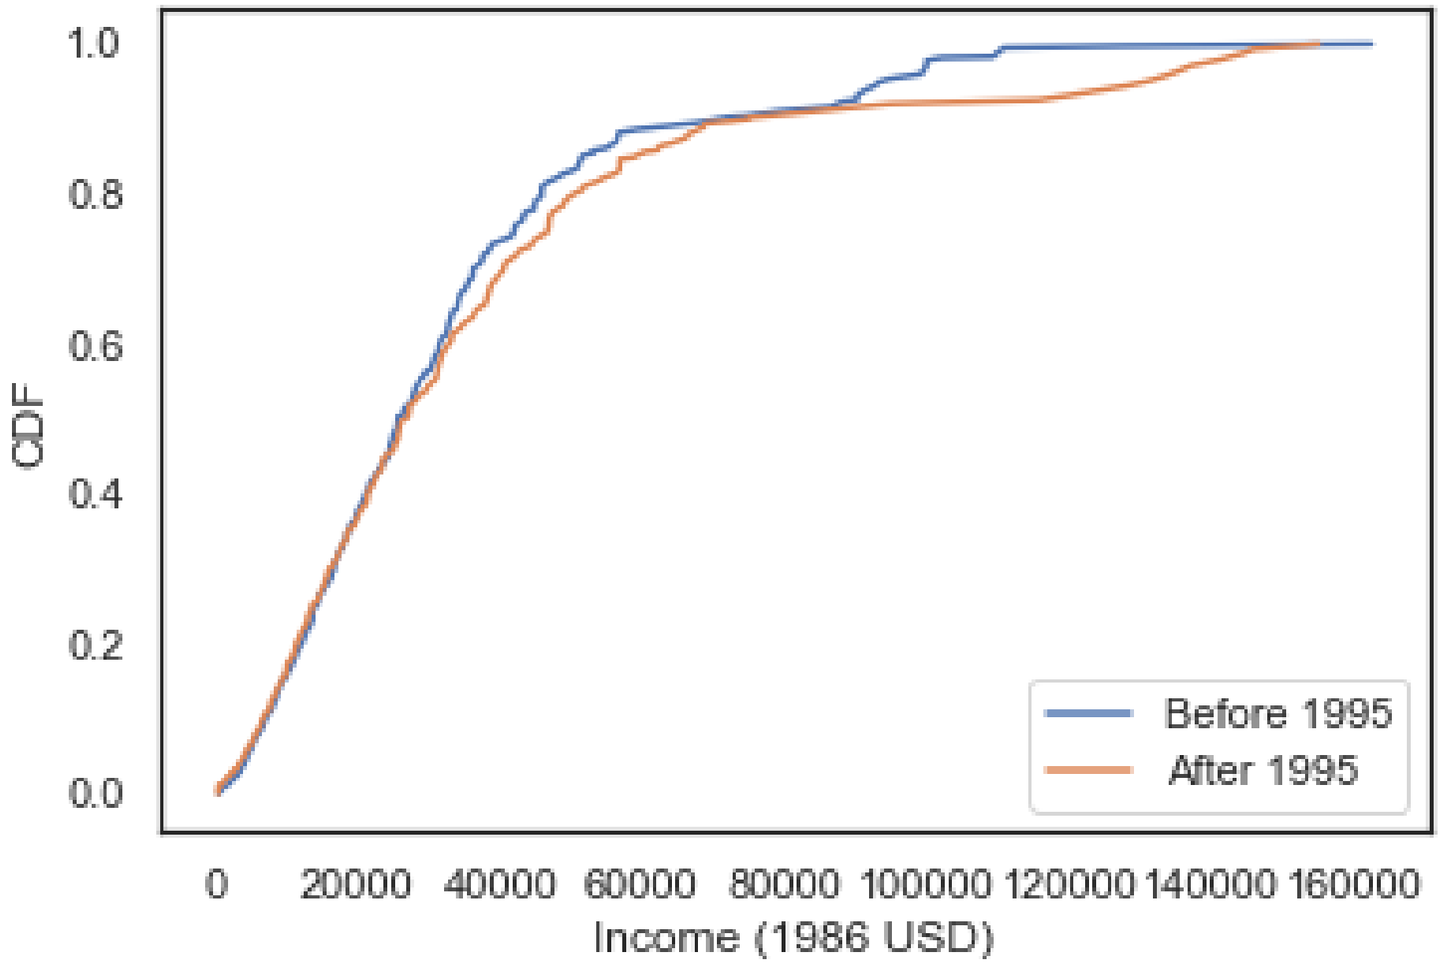
1.0
0.8
0.6
0.4
0.2
0.0
CDF
0
20000
40000
60000
80000
100000
120000
140000
160000
Income (1986 USD)
Before 1995
After 1995

Let's practice!

EXPLORATORY DATA ANALYSIS IN PYTHON

Modeling distributions

EXPLORATORY DATA ANALYSIS IN PYTHON

Allen Downey
Professor, Olin College

datacamp

The normal distribution

```
sample = np.random.normal(size=1000)
Cdf(sample).plot()
```

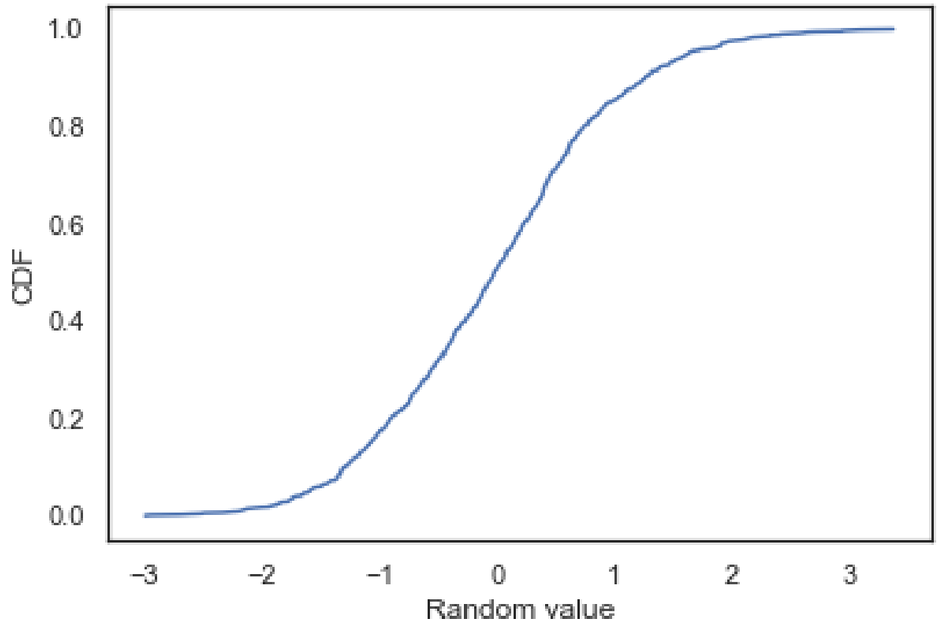

The normal CDF

```
from scipy.stats import norm
```

```
xs = np.linspace(-3, 3)
ys = norm(0, 1).cdf(xs)
```

```
plt.plot(xs, ys, color='gray')
```

```
Cdf(sample).plot()
```

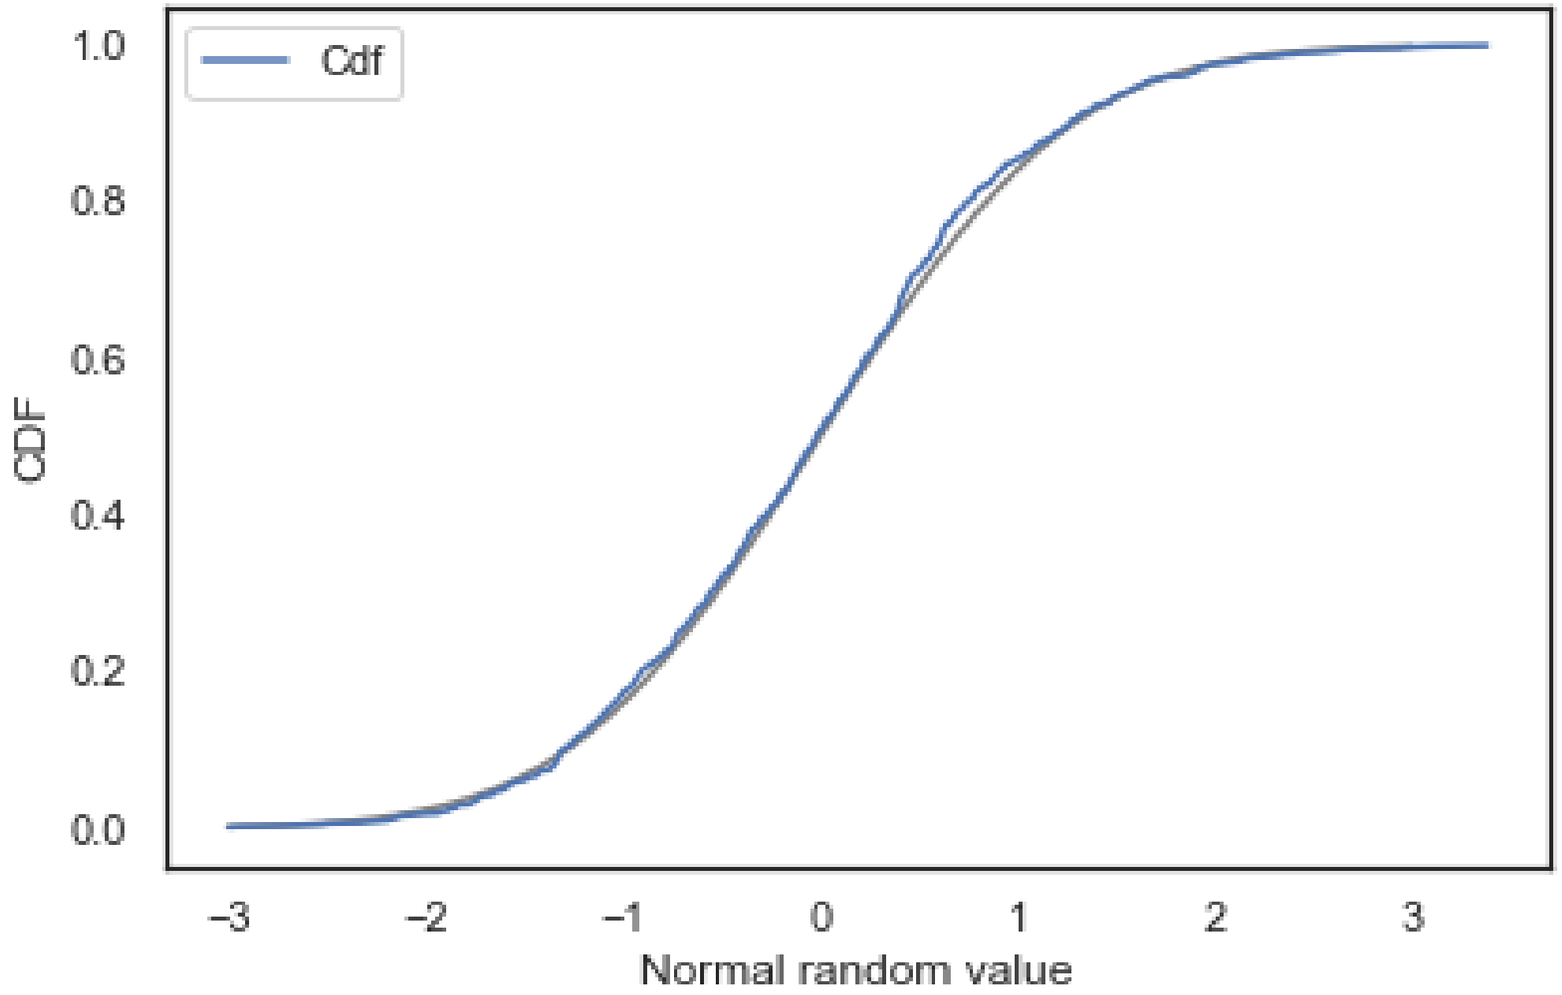
Cdf
1.0
0.8
0.6
0.4
0.2
0.0
CDF
−3
−2
−1
0
1
2
3
Normal random value

The bell curve

```
xs = np.linspace(-3, 3)
ys = norm(0,1).pdf(xs)
plt.plot(xs, ys, color='gray')
```

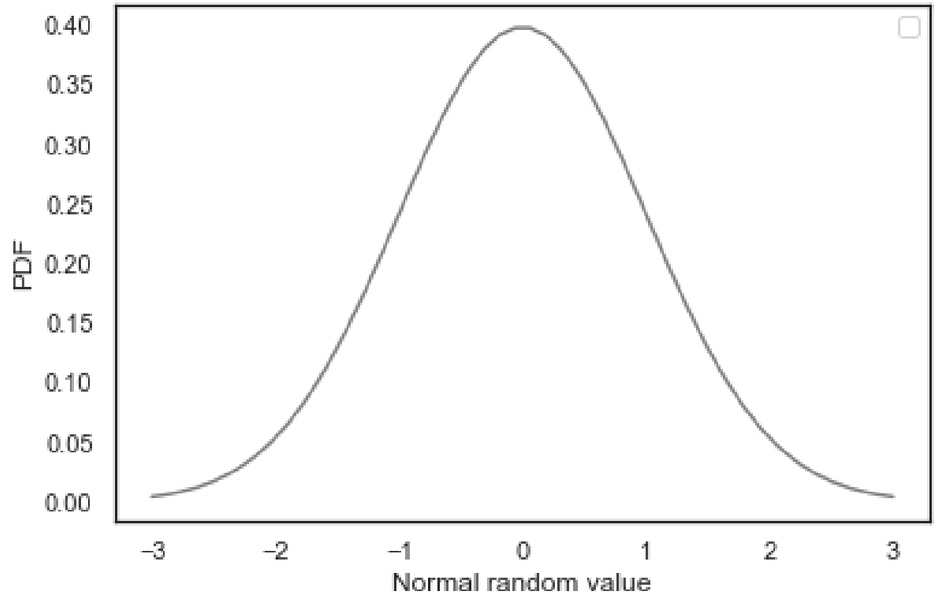

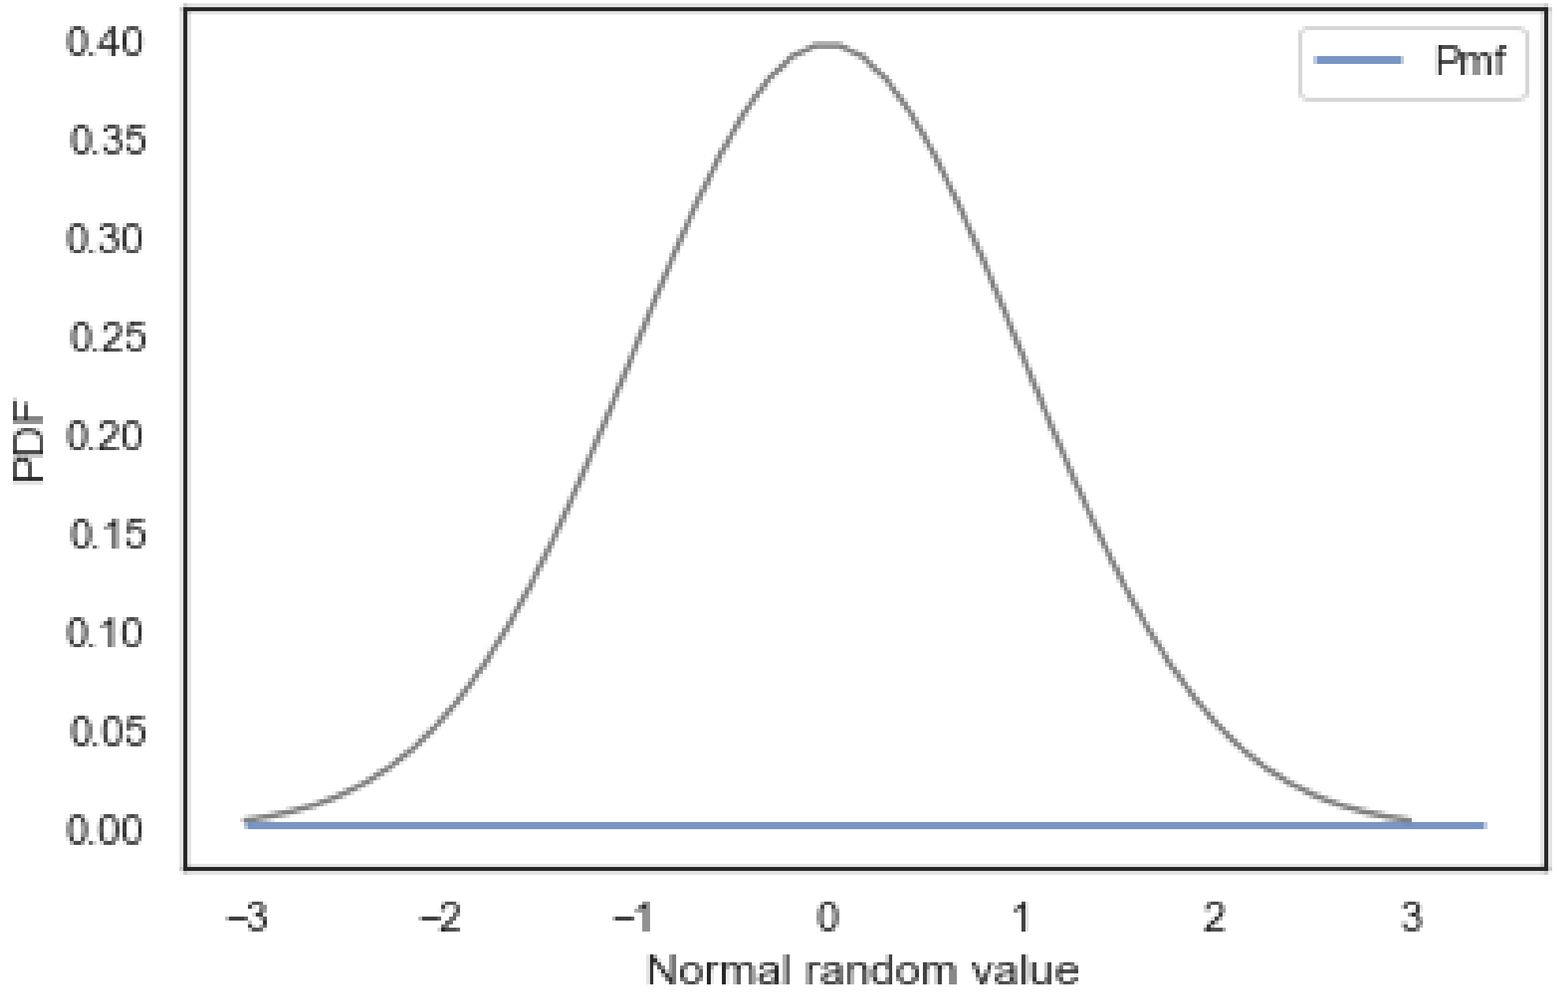
Pmf
0.40
0.35
0.30
0.25
0.20
0.15
0.10
0.05
0.00
PDF
−3
−2
−1
0
1
2
3
Normal random value

KDE plot

```
import seaborn as sns
sns.kdeplot(sample)
```

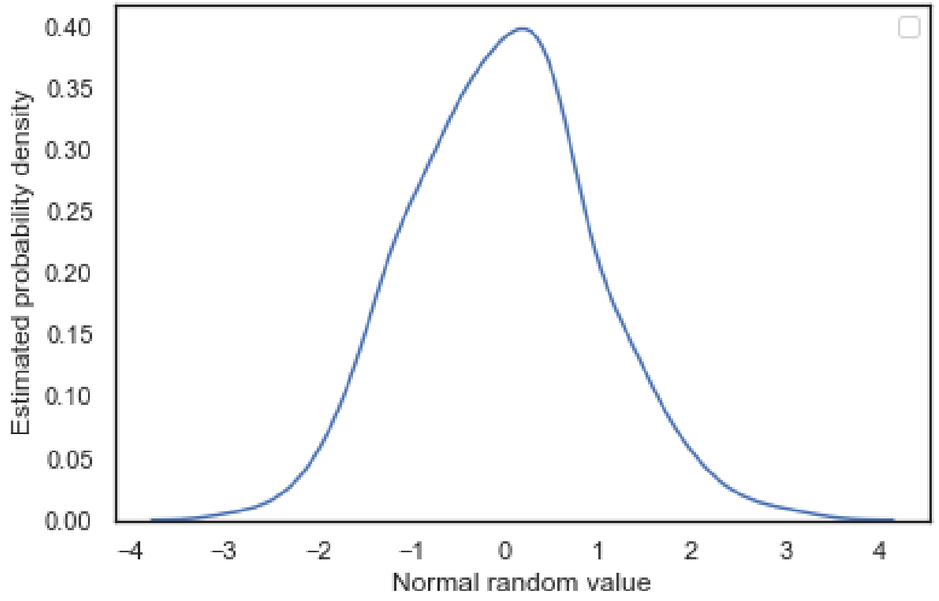

KDE and PDF

```
xs = np.linspace(-3, 3)
ys = norm.pdf(xs)
plt.plot(xs, ys, color='gray')
sns.kdeplot(sample)
```

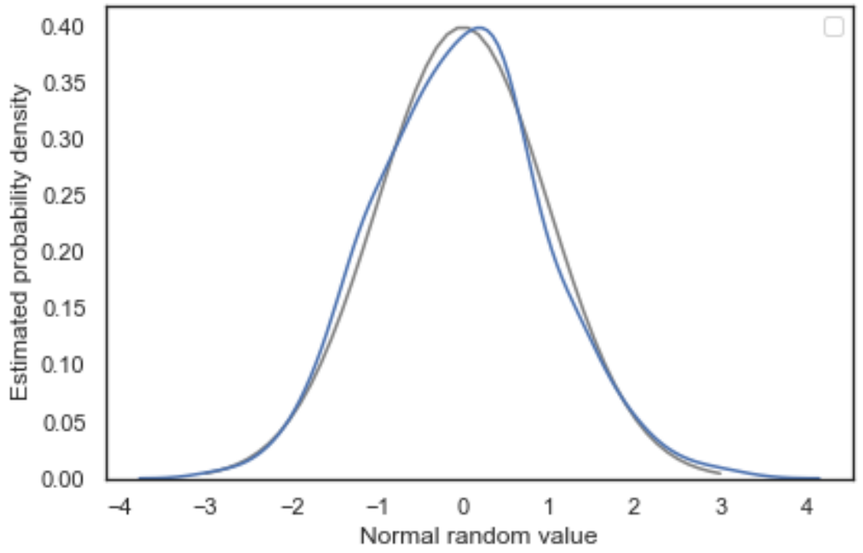

PMF, CDF, KDE

- Use CDFs for exploration.
- Use PMFs if there are a small number of unique values.
- Use KDE if there are a lot of values.

Let's practice!

EXPLORATORY DATA ANALYSIS IN PYTHON