

# Regression review

EXTREME GRADIENT BOOSTING WITH XGBOOST




**Sergey Fogelson**

VP of Analytics, Viacom

# Regression basics

- Outcome is real-valued



# Common regression metrics

- Root mean squared error (RMSE)
- Mean absolute error (MAE)

# Computing RMSE

Actual	Predicted
10	20
3	8
6	1

# Computing RMSE

Actual	Predicted	Error
10	20	-10
3	8	-5
6	1	5

# Computing RMSE

Actual	Predicted	Error	Squared Error
10	20	-10	100
3	8	-5	25
6	1	5	25

- Total Squared Error: 150
- Mean Squared Error: 50
- Root Mean Squared Error: 7.07

# Computing MAE

Actual	Predicted	Error
10	20	-10
3	8	-5
6	1	5

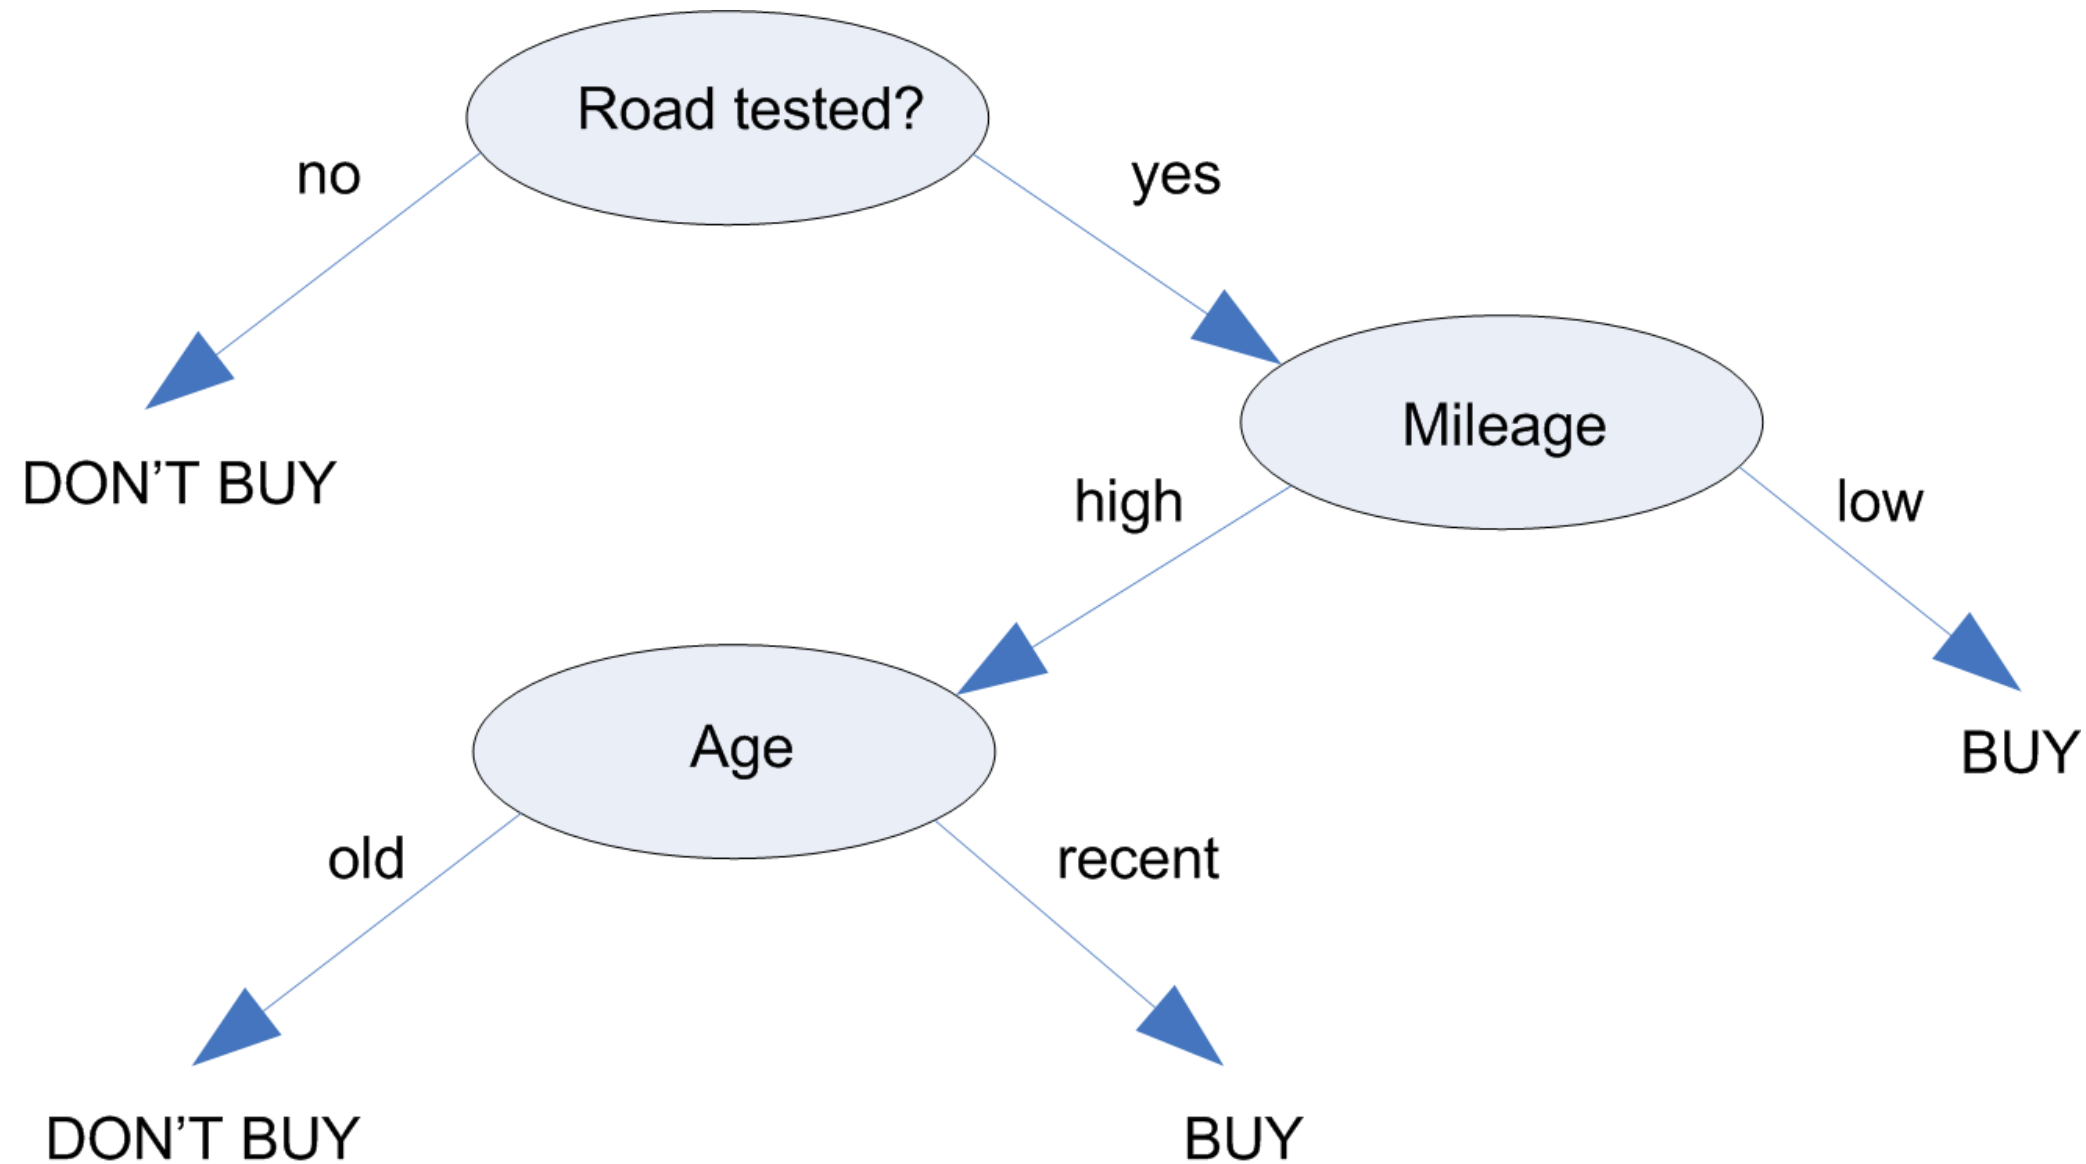
- Total Absolute Error: 20
- Mean Absolute Error: 6.67

# Common regression algorithms

- Linear regression
- Decision trees



# Algorithms for both regression and classification



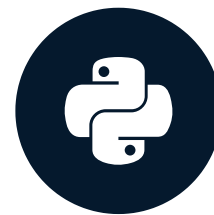
<sup>1</sup> [https://www.ibm.com/support/knowledgecenter/en/SS3RA7\\_15.0.0/com.ibm.spss.modeler.help/nodes\\_treebuilding.htm](https://www.ibm.com/support/knowledgecenter/en/SS3RA7_15.0.0/com.ibm.spss.modeler.help/nodes_treebuilding.htm)

# Let's practice!

EXTREME GRADIENT BOOSTING WITH XGBOOST

# Objective (loss) functions and base learners

EXTREME GRADIENT BOOSTING WITH XGBOOST



**Sergey Fogelson**  
VP of Analytics, Viacom

# Objective Functions and Why We Use Them

- Quantifies how far off a prediction is from the actual result
- Measures the difference between estimated and true values for some collection of data
- Goal: Find the model that yields the minimum value of the loss function

# Common loss functions and XGBoost

- Loss function names in xgboost:
  - reg:linear - use for regression problems
  - reg:logistic - use for classification problems when you want just decision, not probability
  - binary:logistic - use when you want probability rather than just decision

# Base learners and why we need them

- XGBoost involves creating a meta-model that is composed of many individual models that combine to give a final prediction
- Individual models = base learners
- Want base learners that when combined create final prediction that is **non-linear**
- Each base learner should be good at distinguishing or predicting different parts of the dataset
- Two kinds of base learners: tree and linear

# Trees as base learners example: Scikit-learn API

```
import xgboost as xgb
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

boston_data = pd.read_csv("boston_housing.csv")
X, y = boston_data.iloc[:, :-1], boston_data.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=123)

xg_reg = xgb.XGBRegressor(objective='reg:linear', n_estimators=10,
                          seed=123)

xg_reg.fit(X_train, y_train)

preds = xg_reg.predict(X_test)
```

# Trees as base learners example: Scikit-learn API

```
rmse = np.sqrt(mean_squared_error(y_test, preds))  
  
print("RMSE: %f" % (rmse))
```

```
RMSE: 129043.2314
```



# Linear base learners example: learning API only

```
import xgboost as xgb
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

boston_data = pd.read_csv("boston_housing.csv")

X, y = boston_data.iloc[:, :-1], boston_data.iloc[:, -1]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=123)

DM_train = xgb.DMatrix(data=X_train, label=y_train)
DM_test = xgb.DMatrix(data=X_test, label=y_test)

params = {"booster": "gblinear", "objective": "reg:linear"}
xg_reg = xgb.train(params = params, dtrain=DM_train, num_boost_round=10)

preds = xg_reg.predict(DM_test)
```

# Linear base learners example: learning API only

```
rmse = np.sqrt(mean_squared_error(y_test, preds))  
  
print("RMSE: %f" % (rmse))
```

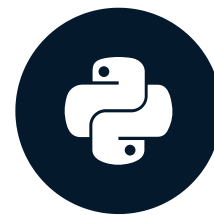
```
RMSE: 124326.24465
```

# Let's get to work!

EXTREME GRADIENT BOOSTING WITH XGBOOST

# Regularization and base learners in XGBoost

EXTREME GRADIENT BOOSTING WITH XGBOOST



**Sergey Fogelson**  
VP of Analytics, Viacom

# Regularization in XGBoost

- Regularization is a control on model complexity
- Want models that are both accurate and as simple as possible
- Regularization parameters in XGBoost:
  - gamma - minimum loss reduction allowed for a split to occur
  - alpha - L1 regularization on leaf weights, larger values mean more regularization
  - lambda - L2 regularization on leaf weights

# L1 regularization in XGBoost example

```
import xgboost as xgb
import pandas as pd
boston_data = pd.read_csv("boston_data.csv")
X,y = boston_data.iloc[:, :-1], boston_data.iloc[:, -1]
boston_dmatrix = xgb.DMatrix(data=X, label=y)
params={"objective":"reg:linear", "max_depth":4}
l1_params = [1, 10, 100]
rmse_l1=[]
for reg in l1_params:
    params["alpha"] = reg
    cv_results = xgb.cv(dtrain=boston_dmatrix, params=params, nfold=4,
                        num_boost_round=10, metrics="rmse", as_pandas=True, seed=123)
    rmse_l1.append(cv_results["test-rmse-mean"].tail(1).values[0])
print("Best rmse as a function of l1:")
print(pd.DataFrame(list(zip(l1_params, rmse_l1)), columns=["l1", "rmse"]))
```

Best rmse as a function of l1:

	l1	rmse
0	1	69572.517742
1	10	73721.967141
2	100	82312.312413

# Base learners in XGBoost

- Linear Base Learner:
  - Sum of linear terms
  - Boosted model is weighted sum of linear models (thus is itself linear)
  - Rarely used
- Tree Base Learner:
  - Decision tree
  - Boosted model is weighted sum of decision trees (nonlinear)
  - Almost exclusively used in XGBoost

# Creating DataFrames from multiple equal-length lists

- ```
pd.DataFrame(list(zip(list1, list2)), columns=["list1", "list2"]))
```
- `zip` creates a `generator` of parallel values:
  - `zip([1, 2, 3], ["a", "b", "c"]) = [1, "a"], [2, "b"], [3, "c"]`
  - `generators` need to be completely instantiated before they can be used in `DataFrame` objects
- `list()` instantiates the full generator and passing that into the `DataFrame` converts the whole expression



# Let's practice!

EXTREME GRADIENT BOOSTING WITH XGBOOST